

Corporate · Legal

Contract Summarisation Agent

CAI Score L4 · Decide — Senior human accountable for every decision

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
TIER	PILLAR	AUTOMATION	HUMAN ROLE	
L4 · Decide	Corporate	≈35% of routine effort	Named accountable decision-maker	

01 — Executive summary

Contract Summarisation Agent is a CAI Score L4 (Decide (supervised)) agent for corporate. Reads contracts and extracts the key risks and obligations, giving legal and business teams a fast, consistent summary. It operates with senior human accountability: a named person owns every decision it supports. Indicatively, it can redirect around 35% of the routine effort this work consumes.

02 — Problem & context

- Lawyers read every contract end-to-end just to find the key risks.
- Business teams wait on legal for a summary they could have had instantly.

03 — Representative use cases


- Produces a consistent risk-and-obligation summary from any contract, MSA, or SLA in minutes.
- Flags items requiring legal review, directing lawyer time to the matters that need it most.
- Provides business teams with a plain-language contract summary without waiting for legal availability.
- Enforces consistent summarisation structure across all contracts, enabling portfolio-level risk analysis.
- Produces the evidence to track summarisation time saved, flagged-item rate, and legal-review efficiency.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (legal).
- 2 Obtain senior sign-off and confirm the named accountable person before go-live.
- 3 Receives the case or work product
- 4 Assembles evidence and flags risk or discrepancy
- 5 Presents structured findings to the named human
- 6 Named human makes the decision; rationale is logged
- 7 Review decision outcomes and the audit trail after each cycle; adjust scope as needed.

05 — Where it fits in the process


The opening step of the contract lifecycle — it produces a consistent risk-and-obligation summary and flags items that need a lawyer eye.



Decision-support loop for a L4 agent – human owns every call.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect evidence assembly	Medium	High	Named human reviews all agent output before any decision is made; the human owns the conclusion.
2	Over-reliance on agent assessment	Low	High	L4 tier structure: agent evidences, human decides. Decision authority cannot be delegated to the agent.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data.
4	Behavioural drift over time	Medium	Medium	Periodic review against case outcomes detects drift; scope and prompts are version-controlled.
5	Out-of-policy or edge-case handling	Low	High	Hard scope boundary: the agent flags anything outside policy; senior sign-off is required at go-live.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Reads contracts and extracts the key risks and obligations, giving legal and business teams a fast, consistent summary.

Assumptions (illustrative — replace with your own)	Result
4 FTE engaged · 1840 productive h/FTE/yr · 35% addressable · €75/h blended	≈ 2,576 hours/yr redirected · ≈ €193,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Reduced contract review workload.

09 — Recommendations

- Obtain senior sign-off before go-live — mandatory for L4 deployments.
- Lawyers own interpretation and legal advice; the agent summarises and flags, it does not advise.
- Assign a named legal lead as the accountable reviewer for all summaries.
- Pilot on 10 contracts before scaling; compare agent summaries against manual summaries for accuracy.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A named senior person owns every decision the agent supports. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.