

Corporate · People Team (HR)

Development & Career Agent

CAI Score L2 · Draft — Human reviews & approves all output

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
TIER	PILLAR	AUTOMATION	HUMAN ROLE	
L2 · Draft	Corporate	≈45% of routine effort	Reviews and approves all output	

01 — Executive summary

Development & Career Agent is a CAI Score L2 (Draft) agent for corporate. Provides career-growth guidance to employees, suggesting development paths and resources. It operates with a human review gate: all output is approved before use. Indicatively, it can redirect around 45% of the routine effort this work consumes.

02 — Problem & context

- Career conversations stall because no one has time to map the options.
- Development resources exist but employees never find the relevant ones.

03 — Representative use cases

- Surfaces relevant development paths and resources for each employee based on their role and stated goals.
- Reduces the preparation burden on managers: the agent provides a structured basis for the development conversation.
- Provides employees with on-demand access to development guidance between formal review cycles.
- Increases the relevance of development resources by personalising suggestions rather than providing a generic catalogue.
- Produces the evidence to track engagement with development resources, conversation frequency, and skill-gap coverage.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (people team (hr)).
- 2 Confirm scope, the named reviewer, and the approval the tier requires.
- 3 Receives the task trigger or request
- 4 Produces the draft or analysis output
- 5 Human reviews and approves before use
- 6 Logs every output for audit
- 7 Review quality and acceptance rate after each cycle; tune scope as you learn.

05 — Where it fits in the process


Sits alongside development conversations — offering personalised growth suggestions and resources that support, not replace, the manager-employee discussion.



Drafting loop for a L2 agent – human approves all output.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete draft	Low	Medium	Draft is reviewed and approved by a human before any use; errors are caught at the review gate.
2	Over-reliance / reduced diligence	Low	Medium	Tier design keeps a human in the approval seat; acceptance-rate KPIs are tracked each cycle.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data. Access stays under your existing controls.
4	Scope creep or misuse	Low	Low	Scope is bounded and documented; the agent produces output only — it takes no autonomous action.
5	Quality drift over time	Low	Low	Periodic sampling and KPI review detects drift; prompts are version-controlled and change-managed.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Provides career-growth guidance to employees, suggesting development paths and resources.

Assumptions (illustrative — replace with your own)	Result
3 FTE engaged · 1840 productive h/FTE/yr · 45% addressable · €50/h blended	≈ 2,484 hours/yr redirected · ≈ €124,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Improved employee engagement.

09 — Recommendations

- Load the approved development framework, career paths, and resource catalogue into the agent.
- Managers and employees make all development decisions — the agent suggests, it does not prescribe.
- Review resource relevance and engagement rates quarterly; update the catalogue as the business evolves.
- Baseline development conversation frequency and employee engagement scores before go-live.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person reviews and approves all output before it is used. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.