

Corporate · People Team (HR)

# Development & Career Agent

CAI Score L2 · Draft — Human reviews & approves all output

▲ this agent

| L1<br>Assist | L2<br>Draft | L3<br>Operate          | L4<br>Decide                    | L5<br>Autonomous |
|--------------|-------------|------------------------|---------------------------------|------------------|
| TIER         | PILLAR      | AUTOMATION             | HUMAN ROLE                      |                  |
| L2 · Draft   | Corporate   | ≈45% of routine effort | Reviews and approves all output |                  |

## 01 — Executive summary

Development & Career Agent is a CAI Score L2 (Draft) agent for corporate. Provides career-growth guidance to employees, suggesting development paths and resources. It operates with a human review gate: all output is approved before use. Indicatively, it can redirect around 45% of the routine effort this work consumes.

## 02 — Problem & context

- Career conversations stall because no one has time to map the options.
- Development resources exist but employees never find the relevant ones.

## 03 — Representative use cases

- Surfaces relevant development paths and resources for each employee based on their role and stated goals.
- Reduces the preparation burden on managers: the agent provides a structured basis for the development conversation.
- Provides employees with on-demand access to development guidance between formal review cycles.
- Increases the relevance of development resources by personalising suggestions rather than providing a generic catalogue.
- Produces the evidence to track engagement with development resources, conversation frequency, and skill-gap coverage.

## 04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (people team (hr)).
- 2 Confirm scope, the named reviewer, and the approval the tier requires.
- 3 Receives the task trigger or request
- 4 Produces the draft or analysis output
- 5 Human reviews and approves before use
- 6 Logs every output for audit
- 7 Review quality and acceptance rate after each cycle; tune scope as you learn.

## 05 — Where it fits in the process


Sits alongside development conversations — offering personalised growth suggestions and resources that support, not replace, the manager-employee discussion.



Drafting loop for a L2 agent – human approves all output.

## 06 — Risk assessment

| # | RISK                              | LIKELIHOOD | IMPACT | MITIGATION / CONTROL                                                                                              |
|---|-----------------------------------|------------|--------|-------------------------------------------------------------------------------------------------------------------|
| 1 | Incorrect or incomplete draft     | Low        | Medium | Draft is reviewed and approved by a human before any use; errors are caught at the review gate.                   |
| 2 | Over-reliance / reduced diligence | Low        | Medium | Tier design keeps a human in the approval seat; acceptance-rate KPIs are tracked each cycle.                      |
| 3 | Exposure of sensitive data        | Medium     | Medium | Runs inside your own tenant; Colleague AI processes none of your data. Access stays under your existing controls. |
| 4 | Scope creep or misuse             | Low        | Low    | Scope is bounded and documented; the agent produces output only — it takes no autonomous action.                  |
| 5 | Quality drift over time           | Low        | Low    | Periodic sampling and KPI review detects drift; prompts are version-controlled and change-managed.                |



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

| FRAMEWORK     | RELEVANCE                                                                            | HOW THIS AGENT ALIGNS                                                                                                                      |
|---------------|--------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| EU AI Act     | Operational use under human oversight; logging and risk-tiering apply.               | Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.                                    |
| DORA          | Relevant where the function is in scope for financial-sector operational resilience. | Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.                             |
| GDPR          | Applies wherever personal data is in scope of the workflow.                          | Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.                           |
| ISO/IEC 42001 | AI management-system controls: roles, monitoring, review.                            | The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified. |

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Provides career-growth guidance to employees, suggesting development paths and resources.

| Assumptions (illustrative — replace with your own)                         | Result                                                       |
|----------------------------------------------------------------------------|--------------------------------------------------------------|
| 3 FTE engaged · 1840 productive h/FTE/yr · 45% addressable · €50/h blended | ≈ 2,484 hours/yr redirected · ≈ €124,000/yr indicative value |



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

**KPIs to track**

Improved employee engagement.

## 09 — Recommendations

- Load the approved development framework, career paths, and resource catalogue into the agent.
- Managers and employees make all development decisions — the agent suggests, it does not prescribe.
- Review resource relevance and engagement rates quarterly; update the catalogue as the business evolves.
- Baseline development conversation frequency and employee engagement scores before go-live.

## 10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person reviews and approves all output before it is used. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

---

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.