

Operations & Service Delivery · Service Delivery

Escalation Triage Agent

CAI Score L3 · Operate — Bounded autonomy

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
--------------	-------------	---------------	--------------	------------------

TIER	PILLAR	AUTOMATION	HUMAN ROLE
L3 · Operate	Operations & Service Delivery	≈65% of routine effort	Owens exceptions and oversight

01 — Executive summary

Escalation Triage Agent is a CAI Score L3 (Operate) agent for operations & service delivery. Handles inbound escalations, classifies severity, drafts an initial response, and routes to the correct owner so issues resolve faster. It operates with bounded autonomy: a person owns exceptions and retains workflow oversight. Indicatively, it can redirect around 65% of the routine effort this work consumes.

02 — Problem & context

- Escalations sit unclassified while the team works out who owns them.
- First responses are drafted from scratch every time, slowing resolution.

03 — Representative use cases

- Classifies every inbound escalation by severity immediately on receipt — no manual triage delay.
- Routes to the correct owner automatically based on severity, client, and category rules.
- Drafts the initial response for the owning team to review and approve before sending.
- Maintains a real-time escalation log with severity, owner, and status for manager visibility.
- Produces the evidence to track resolution time, first-response speed, and routing accuracy.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (service delivery).
- 2 Confirm scope, the named owner, and the oversight the tier requires.
- 3 Detects the trigger in the workflow
- 4 Handles routine in-policy cases end to end
- 5 Escalates exceptions to the owner
- 6 Logs every action for audit
- 7 Review the audit log and KPI movement after each cycle; tune scope as you learn.

05 — Where it fits in the process


The front door for escalations — it classifies and routes the moment one lands, and hands a drafted first response to the owner.



Operating loop for a L3 agent — bounded autonomy.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete output	Medium	Medium	Routine cases only; exceptions and anything material are escalated to the named owner, with a full audit trail.
2	Over-reliance / reduced human scrutiny	Medium	Medium	Tier oversight keeps a human in the loop; KPIs monitor quality so reliance never outruns control.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data, and access stays under your existing controls.
4	Behavioural drift over time	Medium	Low	Periodic review against KPIs and sampling detects drift; scope and prompts are version-controlled.
5	Out-of-policy or edge-case handling	Medium	Medium	Scope is bounded and documented; anything outside it is escalated, not improvised.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Handles inbound escalations, classifies severity, drafts an initial response, and routes to the correct owner so issues resolve faster.

Assumptions (illustrative — replace with your own)	Result
4 FTE engaged · 1840 productive h/FTE/yr · 65% addressable · €55/h blended	≈ 4,784 hours/yr redirected · ≈ €263,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Faster resolution time · reduced escalation risk · consistent first response.

09 — Recommendations

- Define severity classification rules and routing logic before deployment.
- Assign named owners for each severity tier; review the routing accuracy in the first two weeks.
- The owning team approves the drafted first response before sending — this is an L3 agent.
- Baseline resolution time and first-response speed before go-live so the gain is measurable.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person owns exceptions and retains workflow oversight. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI — proprietary. CAI Score™.