

Operations & Service Delivery · Operations · Controls

Four-Eyes Control Assistant

CAI Score L4 · Decide — Senior human accountable for every decision

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
--------------	-------------	---------------	--------------	------------------

TIER	PILLAR	AUTOMATION	HUMAN ROLE
L4 · Decide	Operations & Service Delivery	≈40% of routine effort	Named accountable decision-maker

01 — Executive summary

Four-Eyes Control Assistant is a CAI Score L4 (Decide (supervised)) agent for operations & service delivery. Supports the second reviewer in a four-eyes process — checks the work, highlights discrepancies, and captures the validation evidence. It operates with senior human accountability: a named person owns every decision it supports. Indicatively, it can redirect around 40% of the routine effort this work consumes.

02 — Problem & context

- Dual-control sign-off stalls waiting for a second reviewer to find time.
- Reviewers eyeball everything manually, so discrepancies slip through under deadline.

03 — Representative use cases

- Performs the structured first-pass review in any dual-control chain — finance, compliance, legal, or operations.
- Highlights discrepancies between the work product and source data so the human reviewer sees a diff, not a blank document.
- Accelerates month-end close dual-control: second-review of adjusted entries before journal posting.
- Supports payment and settlement authorisation: structured check of files against thresholds before release.
- Produces a structured evidence pack for every review, satisfying SOX 404 and equivalent ICFR documentation requirements.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (controls).
- 2 Obtain senior sign-off and confirm the named accountable person before go-live.
- 3 Receives the case or work product
- 4 Assembles evidence and flags risk or discrepancy
- 5 Presents structured findings to the named human
- 6 Named human makes the decision; rationale is logged
- 7 Review decision outcomes and the audit trail after each cycle; adjust scope as needed.

05 — Where it fits in the process

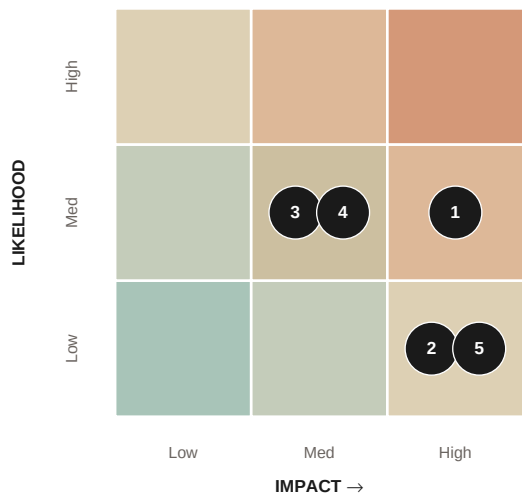
Slots into the dual-control step of any sign-off chain. It does the first-pass check; the human reviewer confirms and remains the named approver.



Decision-support loop for a L4 agent — human owns every call.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect evidence assembly	Medium	High	Named human reviews all agent output before any decision is made; the human owns the conclusion.
2	Over-reliance on agent assessment	Low	High	L4 tier structure: agent evidences, human decides. Decision authority cannot be delegated to the agent.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data.
4	Behavioural drift over time	Medium	Medium	Periodic review against case outcomes detects drift; scope and prompts are version-controlled.
5	Out-of-policy or edge-case handling	Low	High	Hard scope boundary: the agent flags anything outside policy; senior sign-off is required at go-live.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Supports the second reviewer in a four-eyes process — checks the work, highlights discrepancies, and captures the validation evidence.

Assumptions (illustrative — replace with your own)	Result
3 FTE engaged · 1840 productive h/FTE/yr · 40% addressable · €65/h blended	≈ 2,208 hours/yr redirected · ≈ €144,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Faster control cycles · improved auditability · stronger discrepancy detection.

09 — Recommendations

- Require documented senior sign-off (CFO / CRO / General Counsel) before go-live — mandatory for L4.
- Define the named approver role explicitly; the agent will not operate without it.
- Start with a low-volume, well-documented process before extending to higher-stakes chains.
- Engage internal audit at configuration to confirm discrepancy rules meet their evidence standards.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A named senior person owns every decision the agent supports. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.