

Corporate · Reporting

Monthly Reporting Automation Agent

CAI Score L2 · Draft — Human reviews & approves all output

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
TIER	PILLAR	AUTOMATION	HUMAN ROLE	
L2 · Draft	Corporate	≈60% of routine effort	Reviews and approves all output	

01 — Executive summary

Monthly Reporting Automation Agent is a CAI Score L2 (Draft) agent for corporate. Automates recurring monthly reporting, assembling report content and reducing the manual cycle each period. It operates with a human review gate: all output is approved before use. Indicatively, it can redirect around 60% of the routine effort this work consumes.

02 — Problem & context

- An analyst loses a day every month assembling the same report.
- Manual copy-paste between sources introduces errors each cycle.

03 — Representative use cases

- Assembles monthly report content automatically from all connected data sources — no manual copy-paste.
- Eliminates the copy-paste error risk that characterises manual monthly reporting cycles.
- Reduces monthly report preparation from a full analyst day to a review-and-approve task.
- Produces consistent report structure and formatting every cycle, enabling reliable period-on-period comparison.
- Produces the evidence to track preparation time saved, error rate reduction, and reporting consistency.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (reporting).
- 2 Confirm scope, the named reviewer, and the approval the tier requires.
- 3 Receives the task trigger or request
- 4 Produces the draft or analysis output
- 5 Human reviews and approves before use
- 6 Logs every output for audit
- 7 Review quality and acceptance rate after each cycle; tune scope as you learn.

05 — Where it fits in the process


Sits at the monthly reporting step — assembling the report so analysts review rather than rebuild each period.



Drafting loop for a L2 agent – human approves all output.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete draft	Low	Medium	Draft is reviewed and approved by a human before any use; errors are caught at the review gate.
2	Over-reliance / reduced diligence	Low	Medium	Tier design keeps a human in the approval seat; acceptance-rate KPIs are tracked each cycle.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data. Access stays under your existing controls.
4	Scope creep or misuse	Low	Low	Scope is bounded and documented; the agent produces output only — it takes no autonomous action.
5	Quality drift over time	Low	Low	Periodic sampling and KPI review detects drift; prompts are version-controlled and change-managed.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact


FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Automates recurring monthly reporting, assembling report content and reducing the manual cycle each period.

Assumptions (illustrative — replace with your own)	Result
3 FTE engaged · 1840 productive h/FTE/yr · 60% addressable · €55/h blended	≈ 3,312 hours/yr redirected · ≈ €182,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Approx. 1 FTE-day saved per month · reduced error rate.

09 — Recommendations

- Connect the agent to all data sources referenced in the current manual report before deployment.
- An analyst reviews and approves every assembled report before publication — this is an L2 agent.
- Review the first three automated reports against the manual baseline for accuracy.
- Baseline preparation time and error frequency before go-live.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person reviews and approves all output before it is used. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI — proprietary. CAI Score™.