

Corporate · People Team (HR)

Onboarding Agent

CAI Score L2 · Draft — Human reviews & approves all output

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
TIER	PILLAR	AUTOMATION	HUMAN ROLE	
L2 · Draft	Corporate	≈45% of routine effort	Reviews and approves all output	

01 — Executive summary

Onboarding Agent is a CAI Score L2 (Draft) agent for corporate. Guides new joiners through onboarding, answering questions and keeping them on track from day one. It operates with a human review gate: all output is approved before use. Indicatively, it can redirect around 45% of the routine effort this work consumes.

02 — Problem & context

- New joiners wait on HR for answers and lose their first-week momentum.
- Onboarding quality depends on which manager they happened to get.

03 — Representative use cases

- Provides new joiners with on-demand answers to onboarding questions from day one, reducing wait times.
- Guides new joiners through all required onboarding steps in sequence, ensuring nothing is missed.
- Delivers a consistent onboarding experience regardless of manager, team, or location.
- Reduces the volume of routine onboarding queries routed to HR, freeing the People Ops team.
- Produces the evidence to track onboarding completion rates, query volume by topic, and new-joiner satisfaction.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (people team (hr)).
- 2 Confirm scope, the named reviewer, and the approval the tier requires.
- 3 Receives the task trigger or request
- 4 Produces the draft or analysis output
- 5 Human reviews and approves before use
- 6 Logs every output for audit
- 7 Review quality and acceptance rate after each cycle; tune scope as you learn.

05 — Where it fits in the process


Sits with the new joiner from day one — guiding them through the steps and answering common questions so the experience is smooth and consistent.



Drafting loop for a L2 agent – human approves all output.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete draft	Low	Medium	Draft is reviewed and approved by a human before any use; errors are caught at the review gate.
2	Over-reliance / reduced diligence	Low	Medium	Tier design keeps a human in the approval seat; acceptance-rate KPIs are tracked each cycle.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data. Access stays under your existing controls.
4	Scope creep or misuse	Low	Low	Scope is bounded and documented; the agent produces output only — it takes no autonomous action.
5	Quality drift over time	Low	Low	Periodic sampling and KPI review detects drift; prompts are version-controlled and change-managed.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Guides new joiners through onboarding, answering questions and keeping them on track from day one.

Assumptions (illustrative — replace with your own)	Result
4 FTE engaged · 1840 productive h/FTE/yr · 45% addressable · €50/h blended	≈ 3,312 hours/yr redirected · ≈ €166,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Improved new-joiner experience · reduced HR onboarding workload.

09 — Recommendations

- Load all approved onboarding materials, policies, and FAQs into the agent before deployment.
- HR owns the onboarding process; the agent answers from approved content, it does not make policy decisions.
- Review unanswered-query logs weekly for the first month; update content where gaps appear.
- Baseline new-joiner satisfaction scores before go-live.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person reviews and approves all output before it is used. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.