

Data & Infrastructure · Data & Engineering

Power Query Code Generator

CAI Score L2 · Draft — Human reviews & approves all output

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
--------------	-------------	---------------	--------------	------------------

TIER	PILLAR	AUTOMATION	HUMAN ROLE
L2 · Draft	Data & Infrastructure	≈55% of routine effort	Reviews and approves all output

01 — Executive summary

Power Query Code Generator is a CAI Score L2 (Draft) agent for data & infrastructure. Generates Advanced Editor (M) code in Power Query for Excel and Power BI, so analysts do not hand-write transformation code. It operates with a human review gate: all output is approved before use. Indicatively, it can redirect around 55% of the routine effort this work consumes.

02 — Problem & context

- Analysts hand-write M code they half-remember, then debug it for an hour.
- Every new query starts from a blank Advanced Editor.

03 — Representative use cases

- Generates working, commented Power Query M code from a plain-language transformation description.
- Eliminates the blank Advanced Editor: analysts describe what they need and review what they get.
- Accelerates new-analyst onboarding: M fluency is no longer required to produce production-quality queries.
- Documents existing undocumented M code: paste in a query and receive an annotated, explained version.
- Standardises transformation patterns across the BI estate, improving maintainability and peer-review quality.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (data & engineering).
- 2 Confirm scope, the named reviewer, and the approval the tier requires.
- 3 Receives the task trigger or request
- 4 Produces the draft or analysis output
- 5 Human reviews and approves before use
- 6 Logs every output for audit
- 7 Review quality and acceptance rate after each cycle; tune scope as you learn.

05 — Where it fits in the process


Drops into the analyst modelling workflow — describe the transformation, get the M code, paste and check. The saving compounds as the data estate grows.



Drafting loop for a L2 agent — human approves all output.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete draft	Low	Medium	Draft is reviewed and approved by a human before any use; errors are caught at the review gate.
2	Over-reliance / reduced diligence	Low	Medium	Tier design keeps a human in the approval seat; acceptance-rate KPIs are tracked each cycle.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data. Access stays under your existing controls.
4	Scope creep or misuse	Low	Low	Scope is bounded and documented; the agent produces output only — it takes no autonomous action.
5	Quality drift over time	Low	Low	Periodic sampling and KPI review detects drift; prompts are version-controlled and change-managed.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact


FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Generates Advanced Editor (M) code in Power Query for Excel and Power BI, so analysts do not hand-write transformation code.

Assumptions (illustrative — replace with your own)	Result
6 FTE engaged · 1840 productive h/FTE/yr · 55% addressable · €65/h blended	≈ 6,072 hours/yr redirected · ≈ €395,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Approx. 10 min saved per working query · compounds with team scale.

09 — Recommendations

- Deploy to the BI and data engineering team first; extend to finance and operations analysts in phase two.
- Establish a data-input policy: describe schema and logic only; do not include production data values in the prompt.
- Prioritise legacy-code documentation as an early use case — immediate value with zero deployment risk.
- Pair with the Infrastructure Wiki Agent for broader data-team knowledge management coverage.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person reviews and approves all output before it is used. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI — proprietary. CAI Score™.