

Sales & Marketing · Sales

Sales End-to-End Assistant

CAI Score L3 · Operate — Bounded autonomy

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
--------------	-------------	---------------	--------------	------------------

TIER	PILLAR	AUTOMATION	HUMAN ROLE
L3 · Operate	Sales & Marketing	≈60% of routine effort	Owens exceptions and oversight

01 — Executive summary

Sales End-to-End Assistant is a CAI Score L3 (Operate) agent for sales & marketing. Covers the full sales cycle — lead generation, RFP drafting, pipeline analysis and deal-execution support — in one assistant. It operates with bounded autonomy: a person owns exceptions and retains workflow oversight. Indicatively, it can redirect around 60% of the routine effort this work consumes.

02 — Problem & context

- RFPs are drafted from scratch under time pressure, every time.
- Routine cycle admin eats the hours reps should spend selling.

03 — Representative use cases


- Handles routine sales-cycle admin end to end — from lead generation through RFP drafting to pipeline tracking.
- Drafts RFP responses from the knowledge base and prior submissions, cutting authoring time from days to hours.
- Manages pipeline reporting and status updates automatically so managers stop chasing roll-ups.
- Provides deal-execution support: drafts summaries, follow-up communications, and next-step recommendations.
- Produces the evidence to track pipeline velocity, RFP win rate, and admin effort released per rep.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (sales).
- 2 Confirm scope, the named owner, and the oversight the tier requires.
- 3 Detects the trigger in the workflow
- 4 Handles routine in-policy cases end to end
- 5 Escalates exceptions to the owner
- 6 Logs every action for audit
- 7 Review the audit log and KPI movement after each cycle; tune scope as you learn.

05 — Where it fits in the process


Spans the whole cycle as a connective layer, picking up routine work at each stage from first touch through to deal execution.



Operating loop for a L3 agent – bounded autonomy.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete output	Medium	Medium	Routine cases only; exceptions and anything material are escalated to the named owner, with a full audit trail.
2	Over-reliance / reduced human scrutiny	Medium	Medium	Tier oversight keeps a human in the loop; KPIs monitor quality so reliance never outruns control.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data, and access stays under your existing controls.
4	Behavioural drift over time	Medium	Low	Periodic review against KPIs and sampling detects drift; scope and prompts are version-controlled.
5	Out-of-policy or edge-case handling	Medium	Medium	Scope is bounded and documented; anything outside it is escalated, not improvised.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact


FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

Covers the full sales cycle — lead generation, RFP drafting, pipeline analysis and deal-execution support — in one assistant.

Assumptions (illustrative — replace with your own)	Result
5 FTE engaged · 1840 productive h/FTE/yr · 60% addressable · €60/h blended	≈ 5,520 hours/yr redirected · ≈ €331,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Improved revenue performance and higher-quality pipeline.

09 — Recommendations

- Map which routine sales tasks the agent will cover before deployment; configure by stage.
- Sellers own all commitments, pricing decisions, and communications that reach the customer.
- Review the RFP win rate before and after deployment to measure the quality impact.
- Baseline rep admin time as a percentage of selling time before go-live.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person owns exceptions and retains workflow oversight. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.