

Operations & Service Delivery · Service Delivery

Service Delivery Manager Copilot

CAI Score L3 · Operate — Bounded autonomy

▲ this agent

L1 Assist	L2 Draft	L3 Operate	L4 Decide	L5 Autonomous
--------------	-------------	---------------	--------------	------------------

TIER	PILLAR	AUTOMATION	HUMAN ROLE
L3 · Operate	Operations & Service Delivery	≈65% of routine effort	Owens exceptions and oversight

01 — Executive summary

Service Delivery Manager Copilot is a CAI Score L3 (Operate) agent for operations & service delivery. A daily co-pilot for SDMs that pulls together KPI and escalation data per client, flags at-risk accounts, and drafts client-ready updates. It operates with bounded autonomy: a person owns exceptions and retains workflow oversight. Indicatively, it can redirect around 65% of the routine effort this work consumes.

02 — Problem & context

- SDMs start the day blind, hunting KPIs across tools before the first client call.
- An at-risk account only surfaces once the client has already complained.

03 — Representative use cases


- Consolidates KPIs, SLA status, and open escalations per client into a single daily briefing — no manual data hunting.
- Flags at-risk accounts before the client notices, giving SDMs time to proactively manage.
- Drafts client-ready update messages and status summaries for SDM review before sending.
- Provides consistent account visibility across the full client portfolio, not just the noisiest accounts.
- Produces the evidence to track client satisfaction trends, service-credit risk, and escalation patterns.

04 — How to use it

- 1 Connect it in your tenant to the systems this work already touches (service delivery).
- 2 Confirm scope, the named owner, and the oversight the tier requires.
- 3 Detects the trigger in the workflow
- 4 Handles routine in-policy cases end to end
- 5 Escalates exceptions to the owner
- 6 Logs every action for audit
- 7 Review the audit log and KPI movement after each cycle; tune scope as you learn.

05 — Where it fits in the process


Becomes the SDM morning briefing — it consolidates the day KPI and escalation picture per client so the manager walks into the day already ahead of the issues.



Operating loop for a L3 agent – bounded autonomy.

06 — Risk assessment

#	RISK	LIKELIHOOD	IMPACT	MITIGATION / CONTROL
1	Incorrect or incomplete output	Medium	Medium	Routine cases only; exceptions and anything material are escalated to the named owner, with a full audit trail.
2	Over-reliance / reduced human scrutiny	Medium	Medium	Tier oversight keeps a human in the loop; KPIs monitor quality so reliance never outruns control.
3	Exposure of sensitive data	Medium	Medium	Runs inside your own tenant; Colleague AI processes none of your data, and access stays under your existing controls.
4	Behavioural drift over time	Medium	Low	Periodic review against KPIs and sampling detects drift; scope and prompts are version-controlled.
5	Out-of-policy or edge-case handling	Medium	Medium	Scope is bounded and documented; anything outside it is escalated, not improvised.



Residual likelihood × impact after the tier's built-in controls. Numbers map to the table above.

07 — Regulatory & compliance impact

FRAMEWORK	RELEVANCE	HOW THIS AGENT ALIGNS
EU AI Act	Operational use under human oversight; logging and risk-tiering apply.	Documented risk tier, human oversight and event logging per agent — designed to map to the obligations.
DORA	Relevant where the function is in scope for financial-sector operational resilience.	Actions are logged and attributable, supporting resilience and traceability requirements — designed to map to.
GDPR	Applies wherever personal data is in scope of the workflow.	Data stays in your tenant; Colleague AI processes none of it. Purpose limitation and minimisation are preserved.
ISO/IEC 42001	AI management-system controls: roles, monitoring, review.	The agent operates within a documented management system — defined tier, named roles, logging and review. Designed to meet; not certified.

“Designed to map to” / “designed to meet” describe alignment with the CAI Score framework, not external certification. Confirm obligations with your compliance function.

08 — Value & illustrative ROI

A daily co-pilot for SDMs that pulls together KPI and escalation data per client, flags at-risk accounts, and drafts client-ready updates.

Assumptions (illustrative — replace with your own)	Result
5 FTE engaged · 1840 productive h/FTE/yr · 65% addressable · €60/h blended	≈ 5,980 hours/yr redirected · ≈ €359,000/yr indicative value



Figures are an illustrative model to be validated against your own volumes and rates – not a guarantee.

KPIs to track

Improved client satisfaction · reduced service credits · faster issue identification.

09 — Recommendations

- Connect the agent to all data sources the SDM currently checks manually — ticket tools, KPI dashboards, SLA trackers.
- Assign the named exception-owner role to each SDM for their portfolio.
- Review draft communications before sending; calibrate the tone and content with the SDM team in the first two weeks.
- Baseline client satisfaction and service-credit frequency before go-live.

10 — Governance & method

This agent runs in your own Microsoft/Azure tenant; Colleague AI hosts only the control plane (score, policy, audit) and processes none of your data. A person owns exceptions and retains workflow oversight. Every action is time-stamped and attributable. The CAI Score classification is documented and open to challenge.

This dossier is an evaluation document. ROI and automation figures are illustrative models, not commitments. Compliance statements describe the CAI Score framework and are not a certification of your compliance with any law or standard. © Colleague AI – proprietary. CAI Score™.